

Introduction to Pandas

October 30, 2019

1 Pandas

pandas is an open source library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language. In this notebook we will cover the following topics:

- Series
- DataFrame
- Dropping Entries
- Indexing, Selecting, Filtering
- Arithmetic and Data Alignment
- Function Application and Mapping
- Sorting
- Axis Indices with Duplicate Values
- Summarising and Computing Descriptive Statistics
- Cleaning Data
- Input and Output

For help please refer to [The official documentation page](#).

1.1 Imports

```
In [1]: import pandas as pd
import numpy as np
```

1.2 Series

A Series is a one-dimensional array-like object containing an array of data and an associated array of data labels. The data can be any NumPy data type and the labels are the Series' indices.

```
In [2]: # 1) Create a List
myList = [1, 2, 3, -3, 0, 2, 1]
# 2) Convert the list to a series
Series1 = pd.Series(myList)
Series1
```

```
Out[2]: 0    1
        1    2
```

```

2    3
3   -3
4    0
5    2
6    1
dtype: int64

```

Note that each element of the list now has an index when it's converted to a series.
Get the array representation of a Series:

```
In [3]: Series1.values
```

```
Out[3]: array([ 1,  2,  3, -3,  0,  2,  1])
```

Get the index of the Series:

```
In [4]: Series1.index
```

```
Out[4]: RangeIndex(start=0, stop=7, step=1)
```

Index objects are immutable and hold the axis labels and metadata such as names and axis names. Now let's create a series with a custom index:

```
In [5]: Series2 = pd.Series(myList, index=['a', 'John', 'c', 'd', '-1', 'f', 'g'])
Series2
```

```
Out[5]: a      1
        John  2
         c    3
         d   -3
        -1    0
         f    2
         g    1
dtype: int64

```

Get a value from a Series:

```
In [6]: Series2[2]
```

```
Out[6]: 3
```

Verify the index number againsts the index name:

```
In [7]: Series2[2] == Series2['c']
```

```
Out[7]: True
```

Get a set of values from a Series by passing in a list of indices:

```
In [8]: Series2[['a', '-1', 'John']]
```

```
Out[8]: a      1
      -1      0
      John    2
      dtype: int64
```

Get values greater than 1:

```
In [9]: Series2[Series2 > 1]
```

```
Out[9]: John    2
      c      3
      f      2
      dtype: int64
```

Multiply by a scalar:

```
In [10]: Series2 * 5
```

```
Out[10]: a      5
      John    10
      c      15
      d     -15
      -1      0
      f      10
      g      5
      dtype: int64
```

Apply a function

```
In [11]: np.exp(Series2)
```

```
Out[11]: a      2.718282
      John    7.389056
      c     20.085537
      d      0.049787
      -1      1.000000
      f      7.389056
      g      2.718282
      dtype: float64
```

A Series is like a fixed-length, ordered dictionary. We can create a series from dictionaries:

```
In [12]: # 1) Create a dictionary with keys as: A, B, C, and values as: 1, 2, 100
      dict1 = {"A": 1, 'B': 2, 'C': 100}
      # 2) Create a Series from the dictionary
      Series3 = pd.Series(dict1)
      Series3
```

```
Out[12]: A      1
      B      2
      C     100
      dtype: int64
```

Note that the keys have become the indices in the Series.

We can also re-order a Series by passing in an index list (indices which are not found are considered as NaN) when creating from a dictionary:

```
In [13]: index_list = ['C', 'B', 'A', 'D']
         Series4 = pd.Series(dict1, index=index_list)
         Series4
```

```
Out[13]: C    100.0
         B     2.0
         A     1.0
         D    NaN
         dtype: float64
```

We can also check for NaNs:

```
In [14]: Series4.isnull()
```

```
Out[14]: C    False
         B    False
         A    False
         D     True
         dtype: bool
```

Or:

```
In [15]: pd.isnull(Series4)
```

```
Out[15]: C    False
         B    False
         A    False
         D     True
         dtype: bool
```

Series automatically aligns differently indexed data in arithmetic operations:

```
In [16]: Series3
```

```
Out[16]: A     1
         B     2
         C    100
         dtype: int64
```

```
In [17]: Series4
```

```
Out[17]: C    100.0
         B     2.0
         A     1.0
         D    NaN
         dtype: float64
```

```
In [18]: Series3 + Series4
```

```
Out[18]: A      2.0  
         B      4.0  
         C     200.0  
         D      NaN  
         dtype: float64
```

We can also name a Series and its index:

```
In [19]: Series4.name = 'mySeries'  
         Series4.index.name = "myIndex"  
         Series4
```

```
Out[19]: myIndex  
         C      100.0  
         B       2.0  
         A       1.0  
         D      NaN  
         Name: mySeries, dtype: float64
```

We can rename a Series' index in place:

```
In [20]: Series4.index = ["CC", "BB", "AA", "DD"]  
         Series4
```

```
Out[20]: CC      100.0  
         BB       2.0  
         AA       1.0  
         DD      NaN  
         Name: mySeries, dtype: float64
```

1.3 DataFrame

A DataFrame is a tabular data structure containing an ordered collection of columns. Each column can have a different type. DataFrames have both row and column indices. Row and column operations are treated roughly symmetrically. Columns returned when indexing a DataFrame are views of the underlying data, not a copy. To obtain a copy, use the `copy()` method.

Pandas can create DataFrames in different ways (e.g., reading in a file (txt, json, csv), or from a dictionary). Let's start by creating a DataFrame from a dictionary:

```
In [21]: # 1) Create a dictionary  
         dict2 = {'City': ['London', 'London', 'London', 'New York', 'New York'],  
                  'Year': [2015, 2016, 2017, 2016, 2017],  
                  'Population': [8.60, 8.71, 8.79, 8.61, 8.62]}  
  
         # 2) Create a DataFrame from the dictionary  
         df1 = pd.DataFrame(dict2)  
         df1
```

```
Out [21]:
```

	City	Year	Population
0	London	2015	8.60
1	London	2016	8.71
2	London	2017	8.79
3	New York	2016	8.61
4	New York	2017	8.62

Create a DataFrame specifying a sequence of columns:

```
In [22]: df2 = pd.DataFrame(df1, columns=['Year', 'City', 'Population'])
df2
```

```
Out [22]:
```

	Year	City	Population
0	2015	London	8.60
1	2016	London	8.71
2	2017	London	8.79
3	2016	New York	8.61
4	2017	New York	8.62

Like Series, columns that are not present in the data are NaN:

```
In [23]: df3 = pd.DataFrame(df1, columns=['Year', 'City', 'Population', 'Unemployment'])
df3
```

```
Out [23]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	NaN
1	2016	London	8.71	NaN
2	2017	London	8.79	NaN
3	2016	New York	8.61	NaN
4	2017	New York	8.62	NaN

We can retrieve a column by the column name, returning a Series:

```
In [24]: df3['City']
```

```
Out [24]:
```

0	London
1	London
2	London
3	New York
4	New York

Name: City, dtype: object

We can retrieve a column by attribute, returning a Series:

```
In [25]: df3.Year
```

```
Out [25]:
```

0	2015
1	2016
2	2017
3	2016
4	2017

Name: Year, dtype: int64

We can retrieve a row by position:

```
In [26]: df3.iloc[2]
```

```
Out[26]: Year          2017
City          London
Population    8.79
Unemployment   NaN
Name: 2, dtype: object
```

We can update a column by assignment:

```
In [27]: df3['Unemployment'] = np.arange(5)
df3
```

```
Out[27]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	0
1	2016	London	8.71	1
2	2017	London	8.79	2
3	2016	New York	8.61	3
4	2017	New York	8.62	4

We can assign a Series to a column (note if assigning a list or array, the length must match the DataFrame, unlike a Series):

```
In [28]: unemployment = pd.Series([5.9, 6.0, 6.2], index=[2, 3, 4])
df3['Unemployment'] = unemployment
df3
```

```
Out[28]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	NaN
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9
3	2016	New York	8.61	6.0
4	2017	New York	8.62	6.2

We can assign a new column that doesn't exist to any existing column to create a new column (a copy):

```
In [29]: df3['Misc'] = df3['City']
df3
```

```
Out[29]:
```

	Year	City	Population	Unemployment	Misc
0	2015	London	8.60	NaN	London
1	2016	London	8.71	NaN	London
2	2017	London	8.79	5.9	London
3	2016	New York	8.61	6.0	New York
4	2017	New York	8.62	6.2	New York

We can also delete the column:

```
In [30]: del df3['Misc']
df3
```

```
Out[30]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	NaN
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9
3	2016	New York	8.61	6.0
4	2017	New York	8.62	6.2

We can create a DataFrame from a nested dictionary of dicts (the keys in the inner dicts are unioned and sorted to form the index in the result, unless an explicit index is specified):

```
In [31]: population = {'London': {2015:8.6, 2016:8.71, 2017:8.79},
                        'New York': {2016:8.61, 2017:8.62}
                        }
df4 = pd.DataFrame(population)
df4
```

```
Out[31]:
```

	London	New York
2015	8.60	NaN
2016	8.71	8.61
2017	8.79	8.62

We can transpose a DataFrame:

```
In [32]: df4.T
```

```
Out[32]:
```

	2015	2016	2017
London	8.6	8.71	8.79
New York	NaN	8.61	8.62

We can set an index name for the DataFrame:

```
In [33]: df4.index.name = 'year'
df4
```

```
Out[33]:
```

	London	New York
year		
2015	8.60	NaN
2016	8.71	8.61
2017	8.79	8.62

We can also set a name for the DataFrame columns

```
In [34]: df4.columns.name = 'City'
df4
```



```
Out [34]: City  London  New York
          year
          2015      8.60      NaN
          2016      8.71      8.61
          2017      8.79      8.62
```

Return the data contained in a DataFrame as a 2D ndarray:

```
In [35]: df4.values
```

```
Out [35]: array([[8.6 ,  nan],
                 [8.71,  8.61],
                 [8.79,  8.62]])
```

1.4 Dropping Entries

```
In [36]: df3
```

```
Out [36]:   Year      City  Population  Unemployment
0  2015   London         8.60             NaN
1  2016   London         8.71             NaN
2  2017   London         8.79             5.9
3  2016 New York         8.61             6.0
4  2017 New York         8.62             6.2
```

Drop rows from a Series or DataFrame:

```
In [37]: df5 = df3.drop([0])
          df5
```

```
Out [37]:   Year      City  Population  Unemployment
1  2016   London         8.71             NaN
2  2017   London         8.79             5.9
3  2016 New York         8.61             6.0
4  2017 New York         8.62             6.2
```

Drop columns from a DataFrame:

```
In [38]: df6 = df5.drop('Unemployment', axis=1)
          df6
```

```
Out [38]:   Year      City  Population
1  2016   London         8.71
2  2017   London         8.79
3  2016 New York         8.61
4  2017 New York         8.62
```

1.5 Indexing, Selecting, Filtering in DataFrames

```
In [39]: df3
```

```
Out[39]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	NaN
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9
3	2016	New York	8.61	6.0
4	2017	New York	8.62	6.2

Select specified columns from a DataFrame:

```
In [40]: df3[['Population', 'City']]
```

```
Out[40]:
```

	Population	City
0	8.60	London
1	8.71	London
2	8.79	London
3	8.61	New York
4	8.62	New York

Select a slice from a DataFrame:

```
In [41]: df3[:3]
```

```
Out[41]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	NaN
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9

or

```
In [42]: df3.iloc[0:3]
```

```
Out[42]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	NaN
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9

Select from a DataFrame based on a filter:

```
In [43]: df3[df3['Population'] > 8.7]
```

```
Out[43]:
```

	Year	City	Population	Unemployment
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9

or

```
In [44]: df3.loc[df3.Population > 8.7]
```

```
Out [44]:
```

	Year	City	Population	Unemployment
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9

Select a slice of rows from a specific column of a DataFrame:

```
In [45]: df3.loc[0:2, "Population"]
```

```
Out [45]:
```

0	8.60
1	8.71
2	8.79

Name: Population, dtype: float64

1.6 Arithmetic and Data Alignment

Adding DataFrame objects results in the union of index pairs for rows and columns if the pairs are not the same, resulting in NaN for indices that do not overlap:

```
In [46]: np.random.seed(0)
df7 = pd.DataFrame(np.random.rand(9).reshape((3, 3)), columns=['a', 'b', 'c'])
df7
```

```
Out [46]:
```

	a	b	c
0	0.548814	0.715189	0.602763
1	0.544883	0.423655	0.645894
2	0.437587	0.891773	0.963663

```
In [47]: np.random.seed(1)
df8 = pd.DataFrame(np.random.rand(9).reshape((3, 3)), columns=['b', 'c', 'd'])
df8
```

```
Out [47]:
```

	b	c	d
0	0.417022	0.720324	0.000114
1	0.302333	0.146756	0.092339
2	0.186260	0.345561	0.396767

```
In [48]: df7 + df8
```

```
Out [48]:
```

	a	b	c	d
0	NaN	1.132211	1.323088	NaN
1	NaN	0.725987	0.792650	NaN
2	NaN	1.078033	1.309223	NaN

```
In [49]: df9 = df8.add(df7, fill_value=0)
df9
```

```
Out [49]:
```

	a	b	c	d
0	0.548814	1.132211	1.323088	0.000114
1	0.544883	0.725987	0.792650	0.092339
2	0.437587	1.078033	1.309223	0.396767

1.7 Function Application and Mapping

```
In [50]: df10 = df8.sub(df9, fill_value=0)
df10
```

```
Out [50]:
```

	a	b	c	d
0	-0.548814	-0.715189	-0.602763	0.0
1	-0.544883	-0.423655	-0.645894	0.0
2	-0.437587	-0.891773	-0.963663	0.0

```
In [51]: df10 = np.abs(df10)
df10
```

```
Out [51]:
```

	a	b	c	d
0	0.548814	0.715189	0.602763	0.0
1	0.544883	0.423655	0.645894	0.0
2	0.437587	0.891773	0.963663	0.0

Apply a function on 1D arrays to each column:

```
In [52]: myFunc = lambda x: x.max() - x.min()
df10.apply(myFunc)
```

```
Out [52]: a    0.111226
b    0.468118
c    0.360899
d    0.000000
dtype: float64
```

Apply a function on 1D arrays to each row:

```
In [53]: df10.apply(myFunc, axis=1)
```

```
Out [53]: 0    0.715189
1    0.645894
2    0.963663
dtype: float64
```

1.8 Sorting

```
In [54]: df11 = pd.DataFrame(np.arange(12).reshape((3, 4)),
                             index=['three', 'one', 'two'],
                             columns=['c', 'a', 'b', 'd'])
df11
```

```
Out [54]:
```

	c	a	b	d
three	0	1	2	3
one	4	5	6	7
two	8	9	10	11

Sort a DataFrame by its index:

```
In [55]: df11.sort_index()
```

```
Out[55]:
```

	c	a	b	d
one	4	5	6	7
three	0	1	2	3
two	8	9	10	11

Sort a DataFrame by columns in descending order:

```
In [56]: df11.sort_index(axis=1, ascending=False)
```

```
Out[56]:
```

	d	c	b	a
three	3	0	2	1
one	7	4	6	5
two	11	8	10	9

1.9 Axis Indices with Duplicate Values

Labels do not have to be unique in Pandas:

Select DataFrame elements:

```
In [57]: df12 = pd.DataFrame(np.random.randn(5, 4),
                             index=['foo', 'foo', 'bar', 'bar', 'baz'])
df12
```

```
Out[57]:
```

	0	1	2	3
foo	-2.363469	1.135345	-1.017014	0.637362
foo	-0.859907	1.772608	-1.110363	0.181214
bar	0.564345	-0.566510	0.729976	0.372994
bar	0.533811	-0.091973	1.913820	0.330797
baz	1.141943	-1.129595	-0.850052	0.960820

```
In [58]: df12.loc['bar']
```

```
Out[58]:
```

	0	1	2	3
bar	0.564345	-0.566510	0.729976	0.372994
bar	0.533811	-0.091973	1.913820	0.330797

1.10 Summarising and Computing Descriptive Statistics

Unlike NumPy arrays, Pandas descriptive statistics automatically exclude missing data. NaN values are excluded unless the entire row or column is NA.

```
In [59]: df3
```

```
Out[59]:
```

	Year	City	Population	Unemployment
0	2015	London	8.60	NaN
1	2016	London	8.71	NaN
2	2017	London	8.79	5.9
3	2016	New York	8.61	6.0
4	2017	New York	8.62	6.2

```
In [60]: df3.sum()
```

```
Out[60]: Year                10081
City      LondonLondonLondonNew YorkNew York
Population                43.33
Unemployment                18.1
dtype: object
```

Sum over the rows:

```
In [61]: df3.sum(axis=1)
```

```
Out[61]: 0    2023.60
1    2024.71
2    2031.69
3    2030.61
4    2031.82
dtype: float64
```

Account for NaNs:

```
In [62]: df3.sum(axis=1, skipna=False)
```

```
Out[62]: 0      NaN
1      NaN
2    2031.69
3    2030.61
4    2031.82
dtype: float64
```

1.11 Cleaning Data

- Replace
- Drop
- Concatenate

1.11.1 Replace

Replace all occurrences of a string with another string, in place (no copy):

```
In [63]: df1
```

```
Out[63]:
```

	City	Year	Population
0	London	2015	8.60
1	London	2016	8.71
2	London	2017	8.79
3	New York	2016	8.61
4	New York	2017	8.62

```
In [64]: df1.replace('London', 'Lon', inplace=True)
df1
```

```
Out [64]:
```

	City	Year	Population
0	Lon	2015	8.60
1	Lon	2016	8.71
2	Lon	2017	8.79
3	New York	2016	8.61
4	New York	2017	8.62

In a specified column, replace all occurrences of a string with another string, in place (no copy):

```
In [65]: df1.replace({'City' : { 'Lon' : 'London' }}, inplace=True)
df1
```

```
Out [65]:
```

	City	Year	Population
0	London	2015	8.60
1	London	2016	8.71
2	London	2017	8.79
3	New York	2016	8.61
4	New York	2017	8.62

1.11.2 Drop

Drop the 'Population' column and return a copy of the DataFrame:

```
In [66]: df13 = df1.drop('Population', axis=1)
df13
```

```
Out [66]:
```

	City	Year
0	London	2015
1	London	2016
2	London	2017
3	New York	2016
4	New York	2017

1.11.3 Concatenate

Concatenate two DataFrames:

```
In [67]: dict3 = {'City': ['Manchester', 'Manchester', 'Manchester', 'Beijing', 'Beijing'],
                  'Year': [2015, 2016, 2017, 2016, 2017],
                  'Population': [2.72, 2.75, 2.81, 21.01, 20.50]}
df14 = pd.DataFrame(dict3)
df14
```

```
Out [67]:
```

	City	Year	Population
0	Manchester	2015	2.72
1	Manchester	2016	2.75
2	Manchester	2017	2.81
3	Beijing	2016	21.01
4	Beijing	2017	20.50

```
In [68]: df15 = pd.concat([df1, df14], axis=0, sort=False)
df15
```

```
Out [68]:
```

	City	Year	Population
0	London	2015	8.60
1	London	2016	8.71
2	London	2017	8.79
3	New York	2016	8.61
4	New York	2017	8.62
0	Manchester	2015	2.72
1	Manchester	2016	2.75
2	Manchester	2017	2.81
3	Beijing	2016	21.01
4	Beijing	2017	20.50

1.12 Input and Output

- Reading
- Writing

1.12.1 Reading

```
In [69]: data = pd.read_csv("oscar_age_female.csv")
```

```
In [70]: data.head()
```

```
Out [70]:
```

	Index	"Year"	"Age"	"Name"	\
1	1928	22	"Janet Gaynor"	"Seventh Heaven	
2	1929	37	"Mary Pickford"	"Coquette"	
3	1930	28	"Norma Shearer"	"The Divorcee"	\t
4	1931	63	"Marie Dressler"	"Min and Bill"	
5	1932	32	"Helen Hayes"	"The Sin of Madelon Claudet"	\t
				"Movie"	
1	Street Angel and Sunrise: A Song of Two Humans"				
2				NaN	
3				NaN	
4				NaN	
5				NaN	

```
In [71]: data.describe()
```

```
Out [71]:
```

	Index	"Year"
count	89.000000	89.000000
mean	1972.000000	36.123596
std	25.836021	11.745231
min	1928.000000	21.000000
25%	1950.000000	28.000000
50%	1972.000000	33.000000
75%	1994.000000	41.000000
max	2016.000000	80.000000

1.12.2 Writing

```
In [72]: data.to_csv("new.csv")
```